

Rahul Rawat

Working Student, AI and Data Scientist in PhD Programs and Embedded Research

PERSONAL DETAILS

Address	C2 16, 68159 Mannheim, Germany
Phone	015563603340
Email	rahulrawat2r@gmail.com
LinkedIn	linkedin.com/in/rahulrawat2r
GitHub	github.com/rahulrawat20022002
Portfolio	rah-portfolio.pages.dev
Date of birth	20 February 2002
Nationality	Indian, student visa with valid work permit
Availability	Werkstudent 20 hours per week immediately, full time from April 2027

PROFILE

Data Science and Analytics Master student at SRH Heidelberg based in Mannheim, with hands on delivery of agentic AI systems, classic machine learning, and applied research translation into working prototypes. I have shipped a Hybrid RAG Orchestrator on Llama **3.1 8b** via Groq with a custom decision making router and LangChain orchestration, a fairness by design credit scoring system with SHAP driven interpretability, an LLM generated plain language explanation layer, and full regulatory documentation, and a recursive time series pipeline at eRay GmbH benchmarking six models head to head with strict anti leakage evaluation. Proficient in Python, pandas, scikit learn, and comfortable across the modern AI stack from classic machine learning to LLMs, I am the right fit for the Working Student role in PhD Programs and Embedded Research at SAP Walldorf.

PROFESSIONAL EXPERIENCE

Oct 2025 to Mar 2026 Heidelberg	eRay GmbH Data Scientist - During a 6 month eRay GmbH and SRH Heidelberg collaboration to forecast lake water quality across 4 target indicators chlorophyll a, turbidity, pH and dissolved oxygen, built an end to end recursive time series pipeline over a 40 feature space with a per target lag suite lag_1h, lag_24h, lag_3d, lag_7d, lag_roll_mean_24h and lag_roll_std_24h. - Benchmarked 6 candidates Ridge, Gradient Boosting, LightGBM, XGBoost, CatBoost and Prophet with strict tree constraints max_depth 4 and learning_rate 0.05, landed on CatBoost MultiQuantile at alpha 0.05, 0.5 and 0.85, producing asymmetric 80 percent prediction intervals that hug the 0 floor and chop the top 15 percent of summer ghost spikes. - Made the September evaluation defensible with a 3 pass outlier system pH tightened from 0 to 14 down to 7.0 to 9.0, Oct and Nov caps of 15.0 on chlorophyll a and 50.0 on turbidity, and a rolling z-score at $z > 2.5$ over 48 hours, and excluded 5 sparse sensors plus 3 concurrent proxies phycocyanin_abs, phycocyanin_abs_comp and toc, surfacing the honest R squared of 0.86 on dissolved oxygen and 0.81 on pH. - Reconstructed Oct and Nov gaps with IterativeImputer MICE, ran full Memory Buffer recalculation across all 6 lag features, generated a synthetic winter canvas with 4 degree Celsius floor and 0.4 degree diurnal amplitude, then wrapped it all in an orchestrator with gate checks, ecological clips dissolved oxygen 4.0 to 18.0 and pH 6.0 to 9.0 and a 0.003 pH per hour velocity clamp.
------------------------------------	--

Technologies: Python, CatBoost, Prophet, scikit learn, MICE

Aug 2023 to Aug
2024
India

SS Engineers and Contractors

Junior Associate Software Developer

- With SS Engineers and Contractors running internal Data Dashboards, Analytics platforms, and Employee Portals used across the company, tasked with building and maintaining the front end features that day to day teams depended on, contributed React UI components across all three internal products, which stayed in daily use by internal teams across the company throughout the year in role.
- With a client relying on a legacy AngularJS app that had to sit inside an existing module federation setup alongside newer React micro frontends, tasked with porting the client's screens across without breaking the running product, ported around **8 routes** from AngularJS to React inside the module federation shell over **4 months** of incremental releases, shipped with no production incidents during rollout.

Technologies: *React, module federation, Playwright, AngularJS, HTML5, CSS3*

Feb 2023 to July
2023
India

SS Engineers and Contractors

Front End Developer Intern

- During a six month internship at SS Engineers and Contractors, tasked with learning the codebase and then taking on small UI work across the internal Data Dashboards and Employee Portals under senior review, paired closely with senior developers to walk the codebase and then shipped focused UI components including charts, filters, and profile pages, iterating each one on code review feedback until it landed.

Technologies: *React, HTML5, CSS3, Git, code review workflow*

PERSONAL PROJECTS

2025

Multi-Agent RAG with LLM-as-Judge and Multilingual EN and DE Support

- For an English only hybrid BM25 plus dense RAG policy analysis system over a **14 document** policy corpus that could not serve German speakers, tasked with adding multilingual support without duplicating the corpus, migrated embeddings and retrieval to a paraphrase multilingual MiniLM L12 v2 shared vector space so a German query retrieves English sources and is answered in German end to end.
- With every agent in the graph re running language detection and producing inconsistent output languages, tasked with a single source of truth, built a LanguageAgent that centralises seeded language detection with a confidence floor, propagates the output language directive to every downstream agent, and routes preprocessing to language matched spaCy pipelines with a blank multilingual fallback and per chunk language recording so retrieval and evaluation can slice by language.
- To measure answer quality without hand grading, implemented an LLM as Judge JudgeAgent scoring answers **1 to 5** on **5 dimensions** (groundedness, relevance, completeness, citation quality, language quality) in JSON mode at temperature 0, and eliminated self preference bias by running the judge on a different local model Qwen2.5 14B from the generator Mistral 7B with a self_judged flag propagated into every report and a hard failure on missing judge model so a silent fallback to self judging cannot regress unnoticed.

Technologies: *Python, LangGraph multi agent, Ollama with Mistral 7B and Qwen2.5 14B, Pinecone, paraphrase multilingual MiniLM L12 v2, spaCy, BM25*

2025

CreditIQ: Fairness by Design Credit Scoring

- For an SRH Heidelberg project on regulated credit scoring where a baseline model was failing the EU AI Act and AGG **80 percent** fairness bar, tasked with getting it back into compliance without gutting predictive quality, applied AIF360 mitigation and threshold calibration on a real credit dataset, which raised the Disparate Impact ratio from a failing **0.79** to a compliant **0.88**.
- After single axis age bias was fixed but younger women were still being penalised, tasked with finding and correcting the hidden intersectional bias, used SHAP driven subgroup analysis to expose the pattern and designed a four way age by gender threshold matrix, which corrected it without over correcting into reverse discrimination.
- With a large false negative rate silently rejecting good applicants, tasked with cutting it while keeping overall accuracy defensible, documented the fairness accuracy trade off as a deliberate and regulator defensible decision, which brought the false negative rate down from **44 percent to 16.7 percent** while accuracy held at **75 percent** on the held out test split.

Technologies: *Python, scikit learn, AIF360, SHAP, Streamlit*

EDUCATION

Apr 2025 to
Present

Heidelberg

M.Sc. Data Science and Analytics

SRH University of Applied Sciences Heidelberg, GPA 1.9

2019 to 2023

Greater Noida,
India

Bachelor of Technology in Computer Science

GL Bajaj Institute of Technology and Management, CGPA 7.3 of 10

RESEARCH AND THESIS

2022 to 2023

Greater Noida,
India

Bachelor Thesis, Diabetes Prediction Using Machine Learning

GL Bajaj Institute of Technology and Management, IEEE style paper

- For a Bachelor thesis on diabetes prediction with a small clinical dataset of **768 patients**, tasked with building a defensible model comparison the examiners could audit, built a full end to end machine learning pipeline comparing six classifiers with **10 fold** cross validation and per model confusion matrices, delivering a model comparison that stood up in the thesis defence.
- Spotting biologically impossible zero values in the source data that the original authors had overlooked, tasked with restoring data integrity before any model fit, applied IQR based outlier removal and proper imputation, which lifted the dataset from silently broken to a clean training input for every downstream model.
- With a **65 to 35 class** imbalance that made accuracy a misleading headline metric, tasked with choosing an evaluation that would not hide errors on the minority class, moved the headline metric from accuracy to ROC AUC, which exposed the real error patterns the accuracy score had been masking and gave the thesis an honest performance comparison.
- With the results needing to be publishable in substance rather than just submissible, tasked with writing them up formally, produced an IEEE style paper including an honest limitations section and what to do differently in a follow up study, which the supervisor accepted as publishable in substance.

Technologies: *Python, scikit learn, Pandas, Seaborn*

CERTIFICATIONS

- NVIDIA: Building LLM Applications With Prompt Engineering, issued November 2025.
- AWS Academy Graduate: AWS Academy Cloud Foundations, issued July 2025.
- Google Data Analytics: Foundations, Data, Data, Everywhere, issued April 2025 on Coursera.

ACHIEVEMENTS

- USAII Global AI Hackathon 2026: Finalist at Graduate Level, awarded by the United States Artificial Intelligence Institute for innovation, technical creativity, and applied AI on real world challenges.

LANGUAGES

English

Fluent, written and spoken

German

B1 in progress toward B2

Hindi

Native