

Rahul Rawat

Working Student, AI, Digital Products and Process Automation for Finance and MA

PERSONAL DETAILS

Address: C2 16, 68159 Mannheim, Germany

Phone: 015563603340

Email: rahulrawat2r@gmail.com

LinkedIn: linkedin.com/in/rahulrawat2r

GitHub: github.com/rahulrawat20022002

Portfolio: rah-portfolio.pages.dev

Date of birth: 20 February 2002

Nationality: Indian, student visa with valid work permit

Availability: Werkstudent 20 hours per week immediately, full time from April 2027

PROFILE

Master's student in Data Science and Analytics at SRH Heidelberg, based in Mannheim, with hands on experience building AI powered solutions, dashboards, and internal automation over Python, SQL, and LLM tooling. I built a multi agent RAG system with an LLM as Judge evaluation running locally on Ollama with Mistral **7B** and Qwen**2.5 14B**, and delivered a recursive time series pipeline for four water quality indicators at eRay GmbH using CatBoost MultiQuantile with **80 percent** prediction intervals. Comfortable across Python, SQL, LangGraph, cloud platforms and web app work in React from earlier full time front end experience.

SKILLS

Python, SQL, PySpark, BigQuery, Databricks, Delta Lake, Power BI, Tableau, Looker Studio, LangChain, LangGraph, Ollama, CatBoost, LightGBM, XGBoost, Prophet, MICE, scikit-learn, PyTorch, dbt, Apache Airflow, GCP, AWS, Docker, Git, React, Playwright

PROFESSIONAL EXPERIENCE

eRay GmbH

Data Scientist · Heidelberg · Oct 2025 to Mar 2026

- During a **6 month** eRay GmbH and SRH Heidelberg collaboration to forecast lake water quality across **4 target** indicators chlorophyll a, turbidity, pH and dissolved oxygen, built an end to end recursive time series pipeline over a **40 feature** space with a per target lag suite lag_1h, lag_24h, lag_3d, lag_7d, lag_roll_mean_24h and lag_roll_std_24h.
- Benchmarked **6 candidates** Ridge, Gradient Boosting, LightGBM, XGBoost, CatBoost and Prophet with strict tree constraints max_depth 4 and learning_rate **0.05**, landed on CatBoost MultiQuantile at alpha **0.05**, **0.5** and **0.85**, producing asymmetric **80 percent** prediction intervals that hug the 0 floor and chop the top **15 percent** of summer ghost spikes.
- Made the September evaluation defensible with a **3 pass** outlier system pH tightened from **0 to 14** down to **7.0 to 9.0**, Oct and Nov caps of **15.0** on chlorophyll a and **50.0** on turbidity, and a rolling z-score at $z > 2.5$ over **48 hours**, and excluded 5 sparse sensors plus 3 concurrent proxies phycocyanin_abs, phycocyanin_abs_comp and toc, surfacing the honest R squared of **0.86** on dissolved oxygen and **0.81** on pH.
- Reconstructed Oct and Nov gaps with IterativeImputer MICE, ran full Memory Buffer recalculation across all 6 lag features, generated a synthetic winter canvas with 4 degree Celsius floor and **0.4** degree diurnal amplitude, then wrapped it all in an orchestrator with gate checks, ecological clips dissolved oxygen **4.0 to 18.0** and pH **6.0 to 9.0** and a **0.003** pH per hour velocity clamp.

Technologies: Python, CatBoost, Prophet, scikit learn, MICE

SS Engineers and Contractors

Junior Associate Software Developer · India · Aug 2023 to Aug 2024

- With SS Engineers and Contractors running internal Data Dashboards, Analytics platforms, and Employee Portals used across the company, tasked with building and maintaining the front end features that day to day teams depended on, contributed React UI components across all three internal products, which stayed in daily use by internal teams across the company throughout the year in role.
- With a client relying on a legacy AngularJS app that had to sit inside an existing module federation setup alongside newer React micro frontends, tasked with porting the client's screens across without breaking the running product, ported around **8 routes** from AngularJS to React inside the module federation shell over **4 months** of incremental releases, shipped with no production incidents during rollout.

Technologies: *React, module federation, Playwright, AngularJS, HTML5, CSS3*

SS Engineers and Contractors

Front End Developer Intern · India · Feb 2023 to July 2023

- During a six month internship at SS Engineers and Contractors, tasked with learning the codebase and then taking on small UI work across the internal Data Dashboards and Employee Portals under senior review, paired closely with senior developers to walk the codebase and then shipped focused UI components including charts, filters, and profile pages, iterating each one on code review feedback until it landed.

Technologies: *React, HTML5, CSS3, Git, code review workflow*

EDUCATION

M.Sc. Data Science and Analytics

SRH University of Applied Sciences Heidelberg, GPA 1.9 · Heidelberg · Apr 2025 to Present

Bachelor of Technology in Computer Science

GL Bajaj Institute of Technology and Management, CGPA 7.3 of 10 · Greater Noida, India · 2019 to 2023

PERSONAL PROJECTS

Multi-Agent RAG with LLM-as-Judge and Multilingual EN and DE Support

2025

- For an English only hybrid BM25 plus dense RAG policy analysis system over a **14 document** policy corpus that could not serve German speakers, tasked with adding multilingual support without duplicating the corpus, migrated embeddings and retrieval to a paraphrase multilingual MiniLM L12 v2 shared vector space so a German query retrieves English sources and is answered in German end to end.
- With every agent in the graph re running language detection and producing inconsistent output languages, tasked with a single source of truth, built a LanguageAgent that centralises seeded language detection with a confidence floor, propagates the output language directive to every downstream agent, and routes preprocessing to language matched spaCy pipelines with a blank multilingual fallback and per chunk language recording so retrieval and evaluation can slice by language.
- To measure answer quality without hand grading, implemented an LLM as Judge JudgeAgent scoring answers **1 to 5** on **5 dimensions** (groundedness, relevance, completeness, citation quality, language quality) in JSON mode at temperature 0, and eliminated self preference bias by running the judge on a different local model Qwen2.5 **14B** from the generator Mistral **7B** with a self_judged flag propagated into every report and a hard failure on missing judge model so a silent fallback to self judging cannot regress unnoticed.

Technologies: *Python, LangGraph multi agent, Ollama with Mistral 7B and Qwen2.5 14B, Pinecone, paraphrase multilingual MiniLM L12 v2, spaCy, BM25*

Movie Analytics and ML Pipeline on GCP

2025

- For a personal cloud data project on movie analytics that needed to run unattended on GCP, tasked with the ingestion and processing layer, built an end to end batch pipeline that pulls movie data from a public API into a GCS data lake and processes it through a **3 tier** Bronze Silver Gold medallion architecture in BigQuery on Cloud Run, running on a fully automated Cloud Scheduler trigger with 0 manual interventions required.
- With raw ingested data being unfit for direct analytics, tasked with hardening the Silver layer, applied schema enforcement, safe type casting, deduplication via window functions and genre normalisation into a relational model, which held clean referential integrity across every downstream Gold table.
- Because a classifier trained on post release signals would leak the answer, tasked with a leakage free hit prediction, trained a BigQuery ML classifier that predicts whether a film will be a hit before release, deliberately splitting features into **2 tables** so only pre release signals feed the model, and confirmed leakage free evaluation on the split.

Technologies: GCP BigQuery, Cloud Run, GCS, Cloud Scheduler, BigQuery ML, Python, SQL, Looker Studio

RESEARCH AND THESIS

Bachelor Thesis, Diabetes Prediction Using Machine Learning

GL Bajaj Institute of Technology and Management, IEEE style paper · Greater Noida, India · 2022 to 2023

- For a Bachelor thesis on diabetes prediction with a small clinical dataset of **768 patients**, tasked with building a defensible model comparison the examiners could audit, built a full end to end machine learning pipeline comparing six classifiers with **10 fold** cross validation and per model confusion matrices, delivering a model comparison that stood up in the thesis defence.
- Spotting biologically impossible zero values in the source data that the original authors had overlooked, tasked with restoring data integrity before any model fit, applied IQR based outlier removal and proper imputation, which lifted the dataset from silently broken to a clean training input for every downstream model.
- With a **65 to 35 class** imbalance that made accuracy a misleading headline metric, tasked with choosing an evaluation that would not hide errors on the minority class, moved the headline metric from accuracy to ROC AUC, which exposed the real error patterns the accuracy score had been masking and gave the thesis an honest performance comparison.
- With the results needing to be publishable in substance rather than just submissible, tasked with writing them up formally, produced an IEEE style paper including an honest limitations section and what to do differently in a follow up study, which the supervisor accepted as publishable in substance.

Technologies: Python, scikit learn, Pandas, Seaborn

CERTIFICATIONS

- NVIDIA: Building LLM Applications With Prompt Engineering, issued November 2025.
- AWS Academy Graduate: AWS Academy Cloud Foundations, issued July 2025.
- Google Data Analytics: Foundations, Data, Data, Everywhere, issued April 2025 on Coursera.

ACHIEVEMENTS

- USAII Global AI Hackathon 2026: Finalist at Graduate Level, awarded by the United States Artificial Intelligence Institute for innovation, technical creativity, and applied AI on real world challenges.

LANGUAGES

English: Fluent, written and spoken

German: B1 in progress toward B2

Hindi: Native